

HUDM6026 Final Project

Chenguang Pan & Seng Lei

April 28, 2023

1.0 Introduction

High School Longitudinal Study of 2009(HSL:09) is a nationally representative, longitudinal study of 23,000+ ninth graders from 944 schools in 2009. It provides comprehensive information about student's background, academic performance in both high school and college, personal attitudes towards study and school, etc. Therefore, this current project uses the HSL:09 open dataset.

This study analyzed a large dataset consisting of 23,503 observations and 9,614 variables to investigate the potential relationship between ninth-grade mathematics foundation and future achievement in STEM fields. To accomplish this objective, a simple linear regression model was employed, which allowed for the estimation of the effect of math proficiency on overall GPA in STEM courses throughout high school. The standardized mathematics assessment of algebraic reasoning, administered during the first semester of grade 9, was utilized to measure students' mathematical abilities at the onset of high school. In turn, the overall GPA in STEM courses was used as a metric of academic performance in STEM subjects throughout the high school years.

After some data cleaning, we kept 199,948 observations for analysis. The simple linear model is:

$$y_i = \beta_0 + \beta_1 \times x_i + e_i,$$

where y_i is the estimated individual outcome for overall STEM GPA, x_i is the student's mathematics assessment score, and e_i is the measurement error.

2.0 Population data descriptions

As a simulated study, we treated cleaned dataset as the population, $N=19948$. The mean and standard deviation for the dependent variable are 2.440 and .934. And 51.250 and 10.031 for the predictor. The correlation coefficient between these two variables is .567.

```
> model_lm <- lm(X3TGPASTEM ~ X1TXMTSCOR, data = hsls_sub)
> summary(model_lm)
```

Call:

```
lm(formula = X3TGPASTEM ~ X1TXMTSCOR, data = hsls_sub)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.82822	-0.50345	0.05364	0.55167	2.70153

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.264846	0.028358	-9.339	<2e-16 ***

```

X1TXMTSCOR    0.052792    0.000543    97.220    <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7693 on 19946 degrees of freedom
Multiple R-squared:  0.3215,    Adjusted R-squared:  0.3215
F-statistic: 9452 on 1 and 19946 DF,  p-value: < 2.2e-16

```

The simple linear regression model presented that the overall model can explain the 32.15% of the variance in outcome, $F(1, 19946) = 9452$, $p < .001$. One score increase in 9th grader's math assessment will be associated with .053 increase in overall STEM GPA and this relation is statistically significant, $\beta_1 = .053$, $p < .001$. The expected value of overall STEM GPA (i.e., β_0) when student gets zero in math assessment is $-.265$, $p < .001$. The negative GPA does not make any sense, but we ignored this issue and move on the study.

3.0 Writing R functions

```

> dat_gen <- function(size= 500, # sample size
+                       betas,    # a numeric array of betas
+                       iv_mean,  # predictor's mean
+                       iv_var,   # predictor's variance
+                       error_sd){ # residuals' sd
+   # data mainly are generated from a normal distribution ~ N(iv_mean, iv_sd)
+   X <- rnorm(size, mean = iv_mean, sd= sqrt(iv_var))
+   X_aug <- cbind(1, X)
+   # residuals are generated from ~N(0, sd)
+   Error <- rnorm(size, mean=0, sd=error_sd)
+   # based on the parameters to generate the outcomes
+   Y <- X_aug %*% as.matrix(betas) + Error
+   out <- cbind(Y, X)
+   colnames(out) <- c("Y", "X1")
+   return(as.data.frame(out))
+ }

```

The data generation function takes the sample size, regression coefficients, predictors' mean and variance, and the standard deviation of residual as input. It returns a simulated dataset in **dataframe** format with the outcome in the first column and predictor in the second.

```

> reg <- function(ds) {
+   x <- as.matrix(ds[,2])
+   y <- as.matrix(ds[,1])
+   y_cen <- apply(y, 2, function(x) x-mean(x))
+   x_cen <- apply(x, 2, function(x) x-mean(x))
+   # the OLS method
+   b1 <- sum(x_cen*y_cen)/sum(x_cen^2)
+   b0 <- mean(y - x*b1)
+   y_hat <- b0 + x*b1
+   sse <- sum((y-y_hat)^2)
+   sig_sq <- sse/(nrow(x)-2)
+   # the alternative method
+   b1_a <- sum(y_cen/x_cen)/nrow(x)

```

```

+   b0_a <- mean(y - x*b1_a)
+   y_hat_a <- b0_a + x*b1_a
+   sse_a <- sum((y-y_hat_a)^2)
+   sig_sq_a <- sse_a/nrow(x)
+   out_ <- cbind(b0, b1, sig_sq, b0_a, b1_a,sig_sq_a)
+   return(out_)
+ }

```

The estimation function takes the simulated data frame as input, and returns the estimated β_0 , β_1 , residual's variance σ^2 from both least square and alternative methods.

4.0 Monte Carlo Simulation

The basic idea behind the Monte Carlo simulation is that one can draw a large number of random samples from a probability distribution representing the population being studied, and then these samples are used to estimate its statistical properties. This project used Monte Carlo method to draw 1000 random samples with the size of 40 by using the `dat_gen()` function above.

```

> R <- 1000
> set.seed(666)
> # randomly generate 1000 samples
> dat_list <- replicate(n = R,
+                       expr = dat_gen(size = 40,
+                                     betas = c(-0.265,0.053),
+                                     iv_mean = 51.24985, iv_var = 100.6209,
+                                     error_sd = 0.7693),
+                       simplify = FALSE)
> # estimated the simple regression model on each sample
> estimates <- sapply(X = dat_list,
+                     FUN = reg,
+                     simplify = TRUE)
> estimates <- t(estimates)
> colnames(estimates) <- c("b0", "b1", "sig_sq", "b0_a", "b1_a", "sig_sq_a")
>
> # write a function to calculate the estimates
> est_out <- function(esti_mat, size){
+   # to calculate the MSE
+   # first to make a parameter matrix in shape of the estimates matrix
+   theta_m <- matrix(c(-0.265, 0.053,0.7693), nrow(esti_mat),6 , byrow = T)
+   # use the estimates matrix minus the parameter matrix
+   est_cent <- esti_mat-theta_m
+   # use apply to get the mse for each estimates
+   estimates_hat_mean <- round(apply(esti_mat,2,mean),3)
+   estimates_hat_var <- round(apply(esti_mat,2,var),3)
+   estimates_hat_se <- round(apply(esti_mat,2,function(x) sd(x)/sqrt(size)),3)
+   estimates_mse <- round(apply(est_cent,2,function(x) sum(x^2)/size),3)
+   results <- rbind(estimates_hat_mean,estimates_hat_var,
+                   estimates_hat_se,estimates_mse)
+   rownames(results) <- c("Mean", "Variance","SE", "MSE")
+   results_trans <- as.data.frame(t(results))
+   # add a column to calculate the bias
+   par_m <- matrix(c(-0.265, 0.053,0.7693), 6,1)

```

```

+ results_trans$Bias <- results_trans$Mean - par_m
+ results_trans$Parameter <- par_m
+ results_trans <- results_trans[,c(6,1,3,2,5,4)]
+ return(results_trans)
+ }
> (mc_out <- est_out(estimates, R))

```

	Parameter	Mean	SE	Variance	Bias	MSE
b0	-0.2650	-0.265	0.020	0.415	0.0000	0.415
b1	0.0530	0.053	0.000	0.000	0.0000	0.000
sig_sq	0.7693	0.585	0.004	0.017	-0.1843	0.051
b0_a	-0.2650	-5.143	1.887	3558.920	-4.8780	3579.161
b1_a	0.0530	0.146	0.036	1.317	0.0930	1.324
sig_sq_a	0.7693	121.727	29.594	875788.198	120.9577	889543.113

5.0 Bootstrap method

The bootstrap method is a statistical technique that involves the repeated sampling with replacement from the original data to obtain estimates of variability and uncertainty of a statistic of interest. Specifically, the statistic is computed on each of the resamples, and the resulting distribution is used to calculate confidence intervals or conduct hypothesis tests. Since we have a single sample of 40 observations in this context, the bootstrap method is proposed to be applied to the row indices of the original dataset. Each resample would be created based on the shuffled indices of the original data, thus allowing for the generation of multiple estimates of the statistic of interest.

The R function for bootstrap is as follows.

```

> # generate a single dataset
> data_b <- dat_gen(size = 40, betas = c(-0.265, 0.053),
+               iv_mean = 51.24985, iv_var = 100.6209,
+               error_sd = 0.7693)
> # run bootstrapping on this single dataset
> B = 1000
> # shuffle the 1:40 index rather than data_b
> boot_index <- replicate(n=B,
+               expr = sample(1:40, 40, TRUE),
+               simplify = FALSE)
> # use the bootstrapped indices to extract the data
> boot_samp <- list()
> for (i in 1:1000) {
+   boot_unit <- data_b[boot_index[[i]],]
+   boot_samp[[i]] <- boot_unit
+ }
> estimates <- sapply(X = boot_samp,
+               FUN = reg,
+               simplify = TRUE)
> estimates <- t(estimates)
> colnames(estimates) <- c("b0", "b1", "sig_sq", "b0_a", "b1_a", "sig_sq_a")
>
> (bs_out <- est_out(estimates, B))

```

	Parameter	Mean	SE	Variance	Bias	MSE
b0	-0.2650	-0.568	0.019	3.670000e-01	-0.3030	4.580000e-01
b1	0.0530	0.058	0.000	0.000000e+00	0.0050	0.000000e+00
sig_sq	0.7693	0.477	0.004	1.200000e-02	-0.2923	9.800000e-02

b0_a	-0.2650	2.819	7.072	5.000976e+04	3.0840	4.996926e+04
b1_a	0.0530	-0.010	0.139	1.922100e+01	-0.0630	1.920600e+01
sig_sq_a	0.7693	3220.335	2961.135	8.768318e+09	3219.5657	8.769915e+09

6.0 Jackknife method

The jackknife method is a statistical technique to estimate the bias and variance of a statistic by systematically leaving out one observation at a time from the dataset, creating a series of “jackknife samples”. By recalculating the statistic of interest on each of these samples, we can obtain an approximation of the distribution of the statistic, and use it to estimate its bias and standard error. In the present context, we applied the jackknife method to the same 40-observation dataset.

```
> jack_list <- list()
> for (i in 1:40) {
+   # each round drop the ith observation
+   data_loov <- data_b[-i,]
+   # make a list to load each jackknife sample
+   jack_list[[i]] <- data_loov
+ }
>
> estimates <- sapply(X = jack_list,
+                     FUN = reg,
+                     simplify = TRUE)
> estimates <- t(estimates)
> # this is a 40 x 6 matrix
> colnames(estimates) <- c("b0", "b1", "sig_sq", "b0_a", "b1_a", "sig_sq_a")
>
> # since the population parameter is known, we use them directly
> (jn_out<-est_out(estimates, 40))
```

	Parameter	Mean	SE	Variance	Bias	MSE
b0	-0.2650	-0.533	0.016	0.010	-0.2680	0.082
b1	0.0530	0.057	0.000	0.000	0.0040	0.000
sig_sq	0.7693	0.510	0.003	0.000	-0.2593	0.068
b0_a	-0.2650	-4.347	6.444	1661.026	-4.0820	1636.160
b1_a	0.0530	0.133	0.129	0.667	0.0800	0.657
sig_sq_a	0.7693	96.400	87.796	308322.929	95.6307	309760.012

One should notice that the population information is available in this simulated study. Therefore, we chose to not use the plug-in principle to estimate the bias or MSE.

7.0 Comparing the results

Table 1 combines the outcomes from all three methods.

Table 1
Statistics from Monte Carlo, Bootstrap, and Jackknife methods

Method	Estimate	Parameter	Mean	SE	Variance	Bias	MSE
Monte Carlo	$\widehat{\beta}_1$	0.053	0.053	0	0	0	0
	$\widehat{\sigma}^2$	0.7693	0.585	0.004	0.017	-0.1843	0.051
	β_1^a	0.053	0.146	0.036	1.317	0.093	1.324
	σ_a^2	0.7693	121.727	29.594	875788.198	120.9577	889543.113
Bootstrap	$\widehat{\beta}_1$	0.053	0.058	0	0	0.005	0
	$\widehat{\sigma}^2$	0.7693	0.477	0.004	0.012	-0.2923	0.098
	β_1^a	0.053	-0.01	0.139	19.221	-0.063	19.206
	σ_a^2	0.7693	3220.335	2961.135	8768318050	3219.5657	8769915333
Jackknife	$\widehat{\beta}_1$	0.053	0.057	0	0	0.004	0
	$\widehat{\sigma}^2$	0.7693	0.51	0.003	0	-0.2593	0.068
	β_1^a	0.053	0.133	0.129	0.667	0.08	0.657
	σ_a^2	0.7693	96.4	87.796	308322.929	95.6307	309760.012